

Neural Nets

Schwartz

September 5, 2016

Making machines that think

The perceptron – a single layer feedforward neural network – was invented in 1957 at the Cornell Aeronautical Laboratory by Frank Rosenblatt with funding from the United States Office of Naval Research. In a 1958 press conference organized by the US Navy, based on Rosenblatt's statements, The New York Times reported the perceptron to be “the embryo of an electronic computer that [the Navy] expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence.”

Upon his death, the titular Reverend Thomas Bayes (1701–1761) of the renowned theorem left an unpublished manuscript deriving the distribution for the parameter of a binomial distribution. This found its way to Richard Price, who prior to posthumously reading the manuscript at the Royal Society, dutifully edited the manuscript and added an introduction that forms much of the philosophical basis for Bayesian analysis. Mathematical historians have argued that “by modern standards, we should refer to the Bayes-Price rule. Price discovered Bayes' work, recognized its importance, corrected it, contributed to the article, and found a use for it. The modern convention of employing Bayes' name alone is unfair but so entrenched that anything else makes little sense.” Quite unaware of Bayes' work, the French mathematician Pierre-Simon Laplace reproduced and extended Bayes' results in 1774. It was also suggested that the blind English mathematician Nicholas Saunderson discovered the theorem some time before Bayes, but this is disputed.

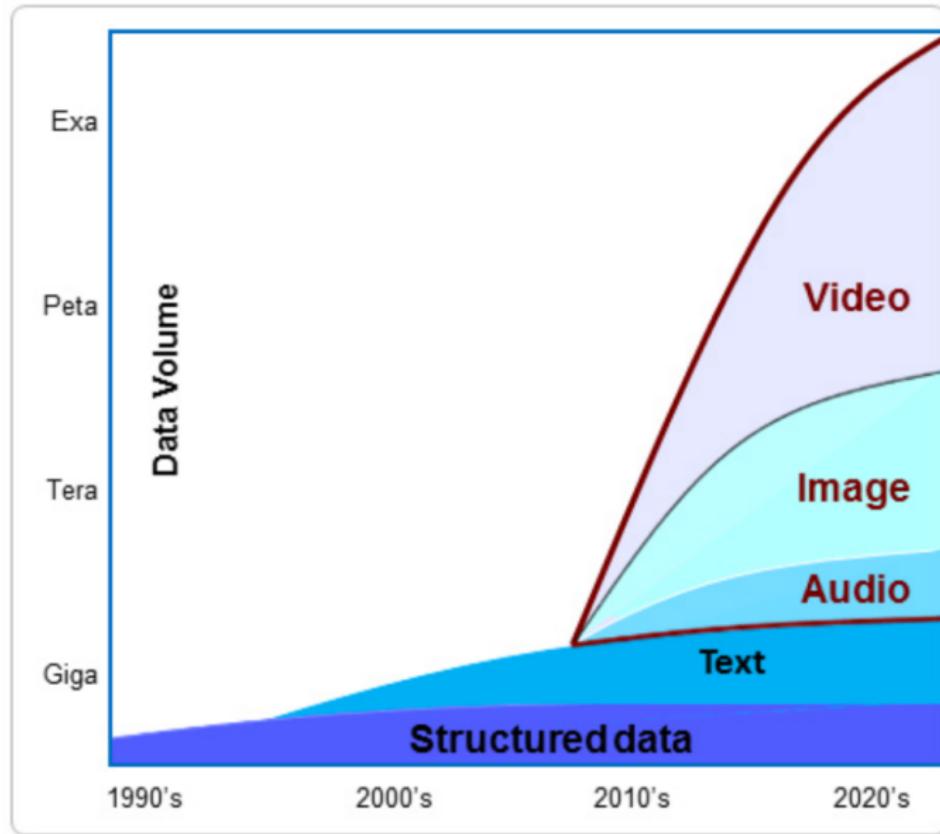
The elegant simplicity and effectiveness of Bayes' theorem for “learning” has prompted some psychologists to ask if the human brain itself might be a Bayesian-reasoning machine. They suggest that the Bayesian capacity to draw strong inferences from sparse data could be crucial to the way the mind perceives the world, plans actions, comprehends and learns language, reasons from correlation to causation, and even understands the goals and beliefs of other minds. The key to successful Bayesian reasoning is not in having an extensive, unbiased sample, which is the eternal worry of frequentists, but rather in having an appropriate “prior”. This prior is an assumption about the way the world works. With the correct prior, even a single piece of data can be used to make meaningful Bayesian predictions. By contrast, frequentism is perhaps not well suited to making decisions on the basis of limited information – which is something that people have to do all the time.

It is thought by cognitive neuroscientists that neocortical development occurs in sequential layers, driven by waves of nerve growth factors which result in a self-organizing system. Modern so-called *deep learning* successors of the perceptron use Bayesian model fitting techniques to analogously sequentially train layer upon layer of neurons to produce unsupervised (self-organizing) classification networks.

Objectives

- ▶ understand image data
- ▶ understand image tools and pipelines
- ▶ understand convolutions
- ▶ understand neural networks
- ▶ understand how they can recapitulate other methodologies
- ▶ understand how activating functions provides power
- ▶ understand how layers provide abstract feature encoding
- ▶ understand how neural networks are trained, i.e.,
“backpropagation is gradient descent is the chain rule”
- ▶ understand that backpropagation doesn’t always work well
- ▶ understand convolutional networks structure and features
- ▶ understand there is new thing called deep learning
that can be viewed as an extension of neural networks

Well that's interesting...



What are images?



How can computers do this??

			
mite black widow cockroach tick starfish	container ship lifeboat amphibian fireboat drilling platform	motor scooter go-kart moped bumper car golfcart	leopard jaguar cheetah snow leopard Egyptian cat
			
grille convertible grille pickup beach wagon fire engine	mushroom agaric mushroom jelly fungus gill fungus dead-man's-fingers	cherry dalmatian grape elderberry ffordshire bulterrier currant	Madagascar cat squirrel monkey spider monkey titi indri howler monkey

How *might* computers be able to do this?

How *might* computers be able to do this?

- ▶ Compress data
- ▶ Keep the search simple
- ▶ Segment image into objects

Image processing *tasks*

- ▶ Read

Image processing tasks

- ▶ Read

`./dog_cat`

`./dog_cat/dog ./dog_cat/cat`

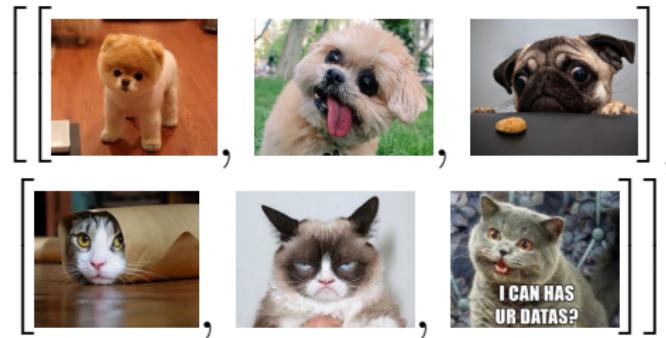
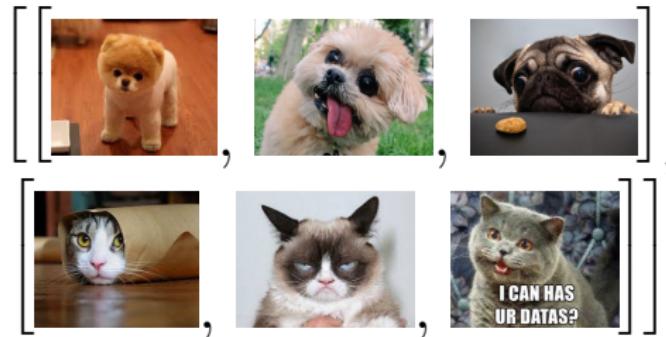


Image processing tasks

- ▶ Read

```
./dog_cat
```

```
./dog_cat/dog ./dog_cat/cat
```



- ▶ `image.shape = (width, height) or (width, height,3)`

Image processing tasks

- ▶ Read

`./dog_cat`

`./dog_cat/dog ./dog_cat/cat`



- ▶ `image.shape = (width, height)` or `(width, height, 3)`
A *tensor* (e.g., a color image) is a matrix of vectors

$$\begin{bmatrix} \textcolor{red}{RGB} & \textcolor{red}{RGB} & \dots & \textcolor{red}{RGB} \\ \textcolor{red}{RGB} & \textcolor{red}{RGB} & \dots & \textcolor{red}{RGB} \\ \vdots & \vdots & \ddots & \vdots \\ \textcolor{red}{RGB} & \textcolor{red}{RGB} & \dots & \textcolor{red}{RGB} \end{bmatrix}$$

Image processing *tasks*

- ▶ Resize

Image processing *tasks*

- ▶ **Resize** to create a uniform feature set

Image processing tasks

- ▶ Resize to create a uniform feature set

Down/Upsampling – *not cropping* – bi-linear interpolation

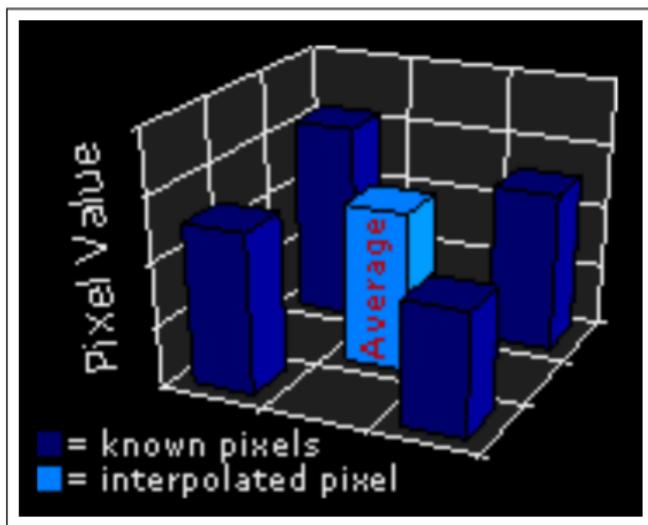
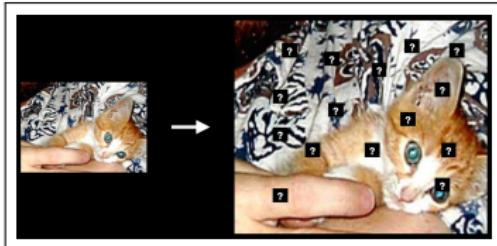
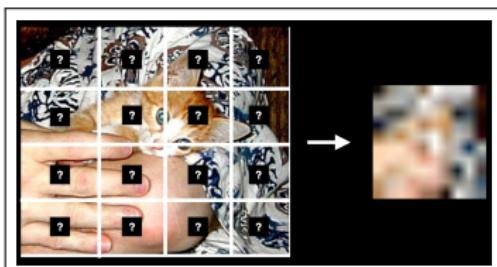
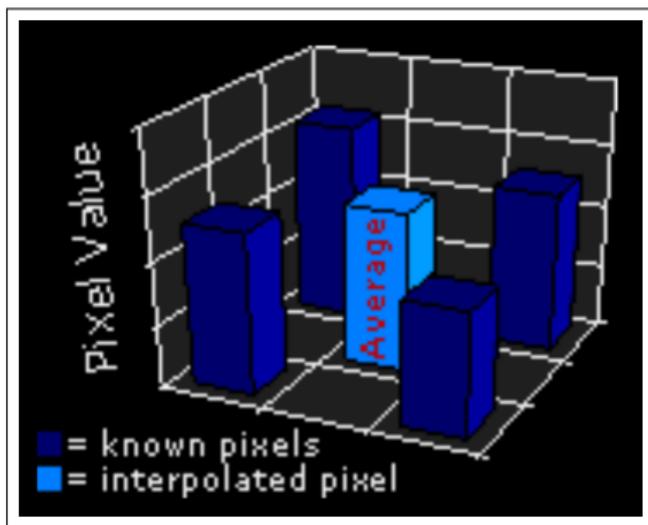
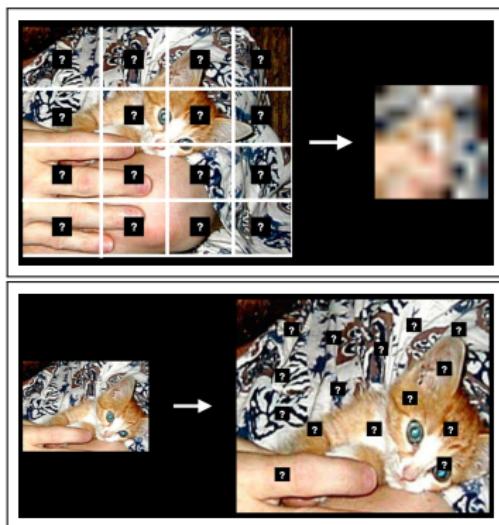


Image processing tasks

- ▶ Resize to create a uniform feature set

Down/Upsampling – *not cropping* – bi-linear interpolation

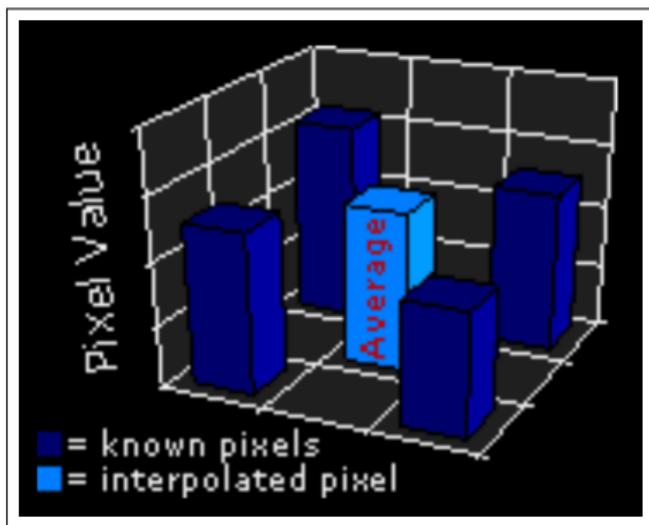
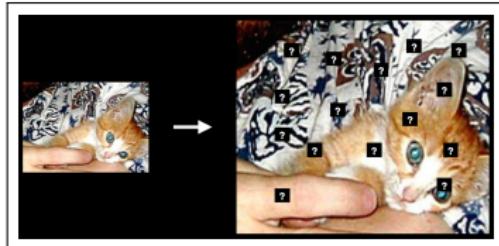
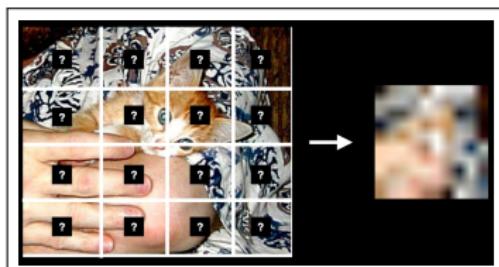


Reduce size/processing time

Image processing tasks

- ▶ Resize to create a uniform feature set

Down/Upsampling – *not cropping* – bi-linear interpolation



Reduce size/processing time
Resolution visually checkable

Image processing tasks

- ▶ Denoise: remove unnecessary details to allow for & better generalization of images to image classes



Image processing tasks – more on denoising

- ▶ Spatial proximity *and* pixel agreement (*bi-lateral*)

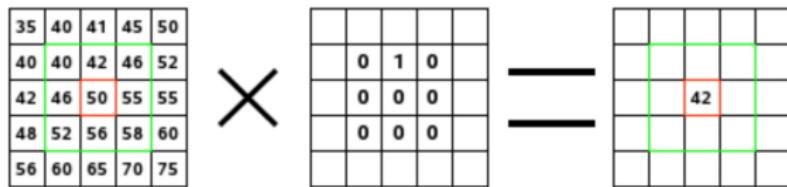


Image processing tasks – more on denoising

- ▶ Spatial proximity *and* pixel agreement (*bi-lateral*)

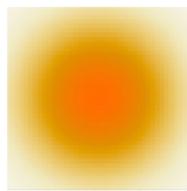
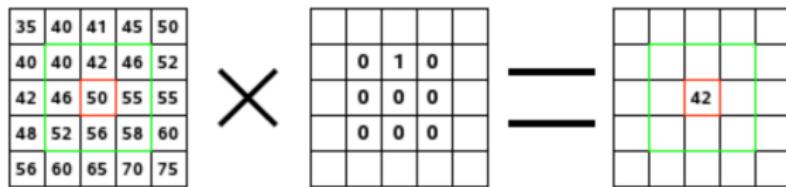


Image processing tasks – more on denoising

- ▶ Spatial proximity *and* pixel agreement (*bi-lateral*)

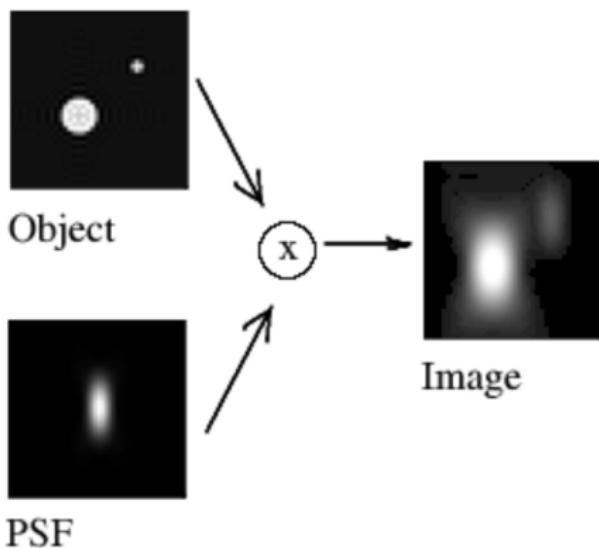
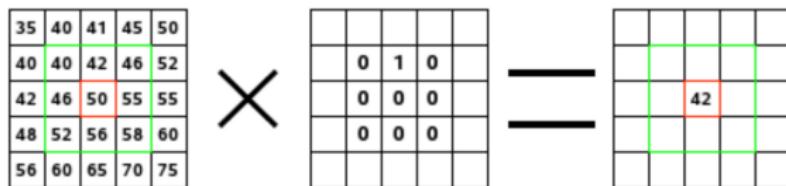


Image processing tasks – more on denoising

- ▶ Spatial proximity *and pixel agreement* (*bi-lateral*)

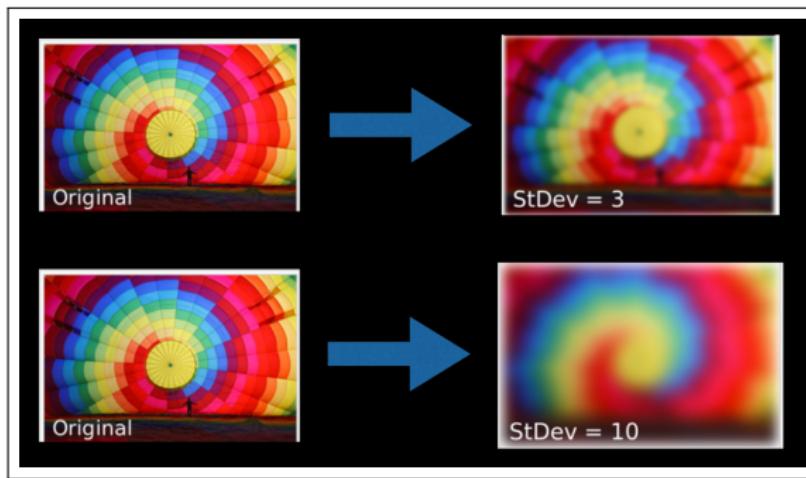
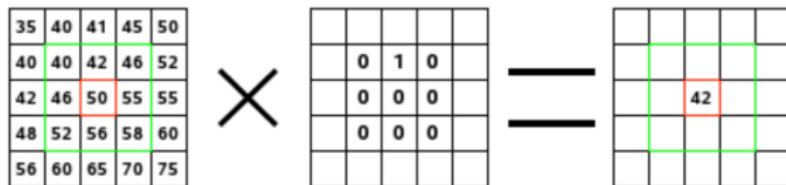


Image processing tasks – more on denoising

- ▶ Variation Minimization: Variation Fidelity tradeoff

$$x_i \rightarrow y_i : \min_y \quad \frac{1}{2} \sum_{Fidelity} (x_i - y_i)^2 + \lambda \sum_{Variation} |y_{i+1} - y_i|$$

Image processing tasks – more on denoising

- ▶ Variation Minimization: Variation Fidelity tradeoff

$$x_i \rightarrow y_i : \min_y \quad \frac{1}{2} \sum_{Fidelity} (x_i - y_i)^2 + \lambda \sum_{Variation} |y_{i+1} - y_i|$$

[What is a “hard” image for this procedure?]

Image processing *tasks*

- ▶ Grayscale (luminescence): Tensor → Matrix

Normalize $0.2125\textcolor{red}{R} + 0.7154\textcolor{green}{G} + 0.0721\textcolor{blue}{B}$

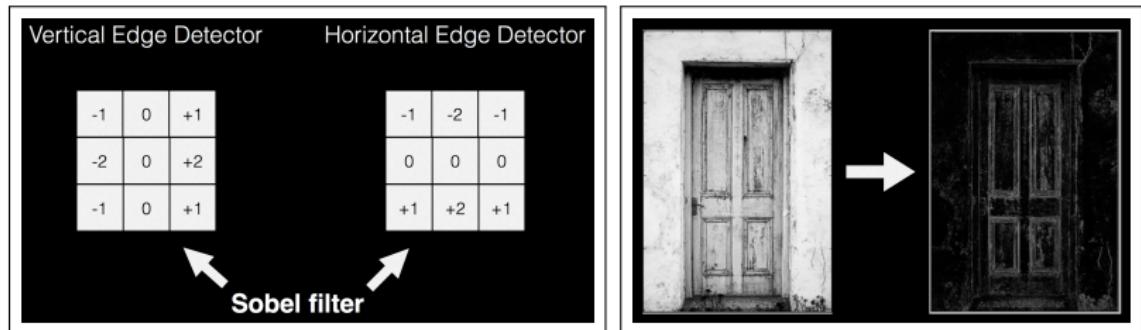
Image processing tasks

- ▶ Grayscale (luminescence): Tensor → Matrix
Normalize $0.2125R + 0.7154G + 0.0721B$
- ▶ Edge detection (grayscale): detects large pixel transitions



Image processing tasks – more on edge-detection

It may be easier to identify objects if we just consider the edges



We may be able to use intensity gradients for feature ascertainment



Image processing *tasks*

- ▶ Image processing libraries
 - ▶ Scikit-image (skimage)
 - ▶ OpenCV (and dependencies)
 - ▶ Python Imaging Library
 - ▶ Pillow (a fork of the above)
 - ▶ etc.

Image Color Structure

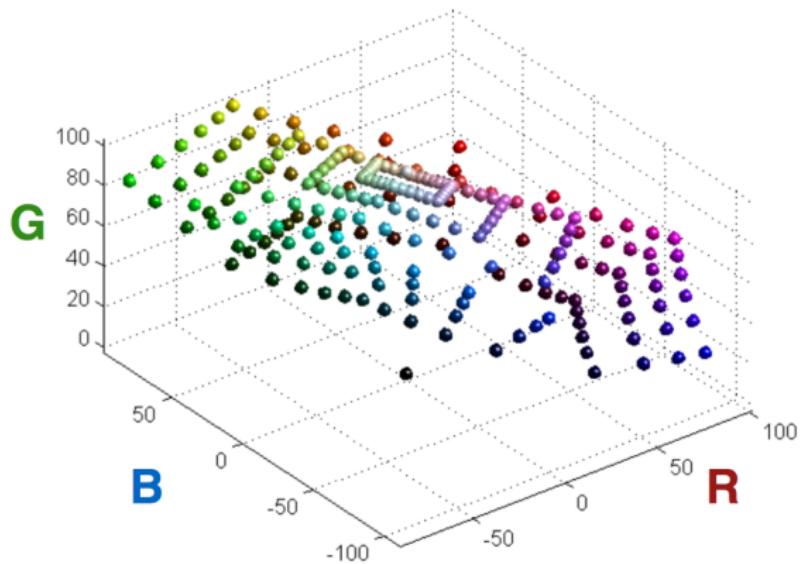


Image Color Structure



Image Featurization

`numpy.ravel()`



`numpy.ravel()`

